


# Graphs

Computer Science E-22  
Harvard Extension School


David G. Sullivan, Ph.D.

## What is a Graph?




- A graph consists of:
  - a set of *vertices* (also known as *nodes*)
  - a set of *edges* (also known as *arcs*), each of which connects a pair of vertices

## Example: A Highway Graph




- Vertices represent cities.
- Edges represent highways.
- This is a *weighted* graph, because it has a *cost* associated with each edge.
  - for this example, the costs denote mileage
- We'll use graph algorithms to answer questions like "What is the shortest route from Portland to Providence?"

## Relationships Among Vertices





- Two vertices are *adjacent* if they are connected by a single edge.
  - ex: c and g are adjacent, but c and i are not
- The collection of vertices that are adjacent to a vertex  $v$  are referred to as  $v$ 's *neighbors*.
  - ex: c's neighbors are a, b, d, f, and g

## Paths in a Graph




- A *path* is a sequence of edges that connects two vertices.
  - ex: the path highlighted above connects c and e
- A graph is *connected* if there is a path between any two vertices.
  - ex: the six vertices at right are part of a graph that is *not* connected
- A graph is *complete* if there is an edge between every pair of vertices.
  - ex: the graph at right is complete



## Directed Graphs


- A *directed* graph has a direction associated with each edge, which is depicted using an arrow:




- Edges in a directed graph are often represented as ordered pairs of the form (start vertex, end vertex).
  - ex: (a, b) is an edge in the graph above, but (b, a) is not.
- A path in a directed graph is a sequence of edges in which the end vertex of edge  $i$  must be the same as the start vertex of edge  $i + 1$ .
  - ex:  $\{ (a, b), (b, e), (e, f) \}$  is a valid path.  
 $\{ (a, b), (c, b), (c, a) \}$  is not.

## Trees vs. Graphs


- A tree is a special type of graph.
  - it is connected and undirected
  - it is *acyclic*: there is no path containing distinct edges that starts and ends at the same vertex
  - we usually single out one of the vertices to be the root of the tree, although graph theory does not require this



a graph that is *not* a tree,  
with one cycle highlighted




a tree using the same nodes




another tree using the same nodes

## Spanning Trees


- A spanning tree is a subset of a connected graph that contains:
  - all of the vertices
  - a subset of the edges that form a tree
- The trees on the previous page were examples of spanning trees for the graph on that page. Here are two others:



the original graph



another spanning tree for this graph




another spanning tree for this graph

## Representing a Graph Using an Adjacency Matrix

- Adjacency matrix = a two-dimensional array that is used to represent the edges and any associated costs
  - $\text{edge}[r][c]$  = the cost of going from vertex  $r$  to vertex  $c$

- Example:

	0	1	2	3
0			54	44
1			39	
2	54	39		83
3	44		83	





- Use a special value to indicate that you can't go from  $r$  to  $c$ .
  - either there's no edge between  $r$  and  $c$ , or it's a directed edge that goes from  $c$  to  $r$
  - this value is shown as a shaded cell in the matrix above
  - we can't use 0, because we may have actual costs of 0
- This representation is good if a graph is *dense* – if it has many edges per vertex – but wastes memory if the graph is *sparse* – if it has few edges per vertex.

## Representing a Graph Using an Adjacency List

- Adjacency list = a list (either an array or linked list) of linked lists that is used to represent the edges and any associated costs

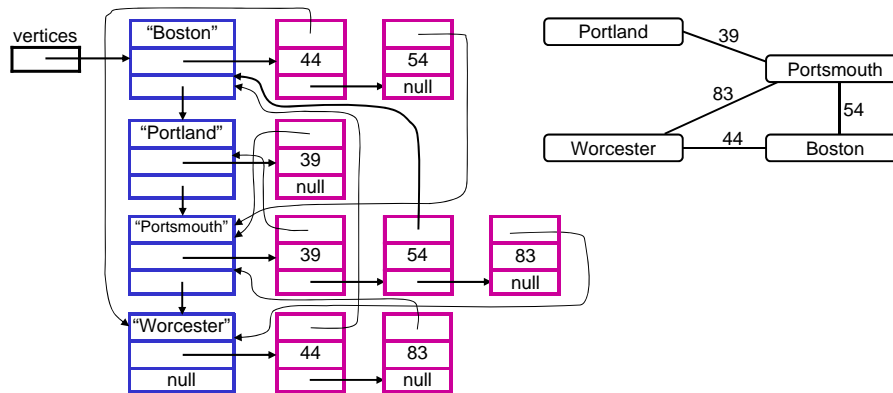
- Example:



- No memory is allocated for non-existent edges, but the references in the linked lists use extra memory.
- This representation is good if a graph is *sparse*, but wastes memory if the graph is *dense*.

## Our Graph Representation

- Use a linked list of linked lists for the adjacency list.
- Example:



vertices is a reference to a linked list of Vertex objects.  
Each Vertex holds a reference to a linked list of Edge objects.  
Each Edge holds a reference to the Vertex that is the end vertex.

## Graph Class

```
public class Graph {
    private class Vertex {
        private String id;
        private Edge edges;           // adjacency list
        private Vertex next;
        private boolean encountered;
        private boolean done;
        private Vertex parent;
        private double cost;
        ...
    }

    private class Edge {
        private Vertex start;
        private Vertex end;
        private double cost;
        private Edge next;
        ...
    }

    private Vertex vertices;
    ...
}
```

*The highlighted fields  
are shown in the diagram  
on the previous page.*

## Traversing a Graph

- Traversing a graph involves starting at some vertex and visiting all of the vertices that can be reached from that vertex.
  - visiting a vertex = processing its data in some way
    - example: print the data
  - if the graph is connected, all of the vertices will be visited
- We will consider two types of traversals:
  - **depth-first**: proceed as far as possible along a given path before backing up
  - **breadth-first**: visit a vertex  
visit all of its neighbors  
visit all unvisited vertices 2 edges away  
visit all unvisited vertices 3 edges away, etc.
- Applications:
  - determining the vertices that can be reached from some vertex
  - state-space search
  - web crawler (vertices = pages, edges = links)

## Depth-First Traversal

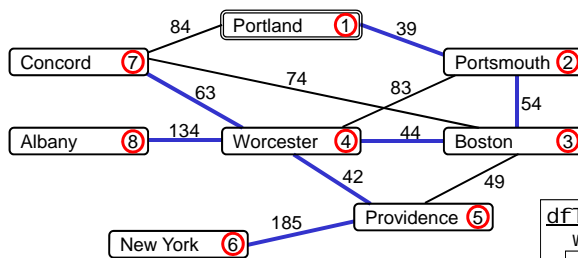
- Visit a vertex, then make recursive calls on all of its yet-to-be-visited neighbors:

```
dfTrav(v, parent)
    visit v and mark it as visited
    v.parent = parent
    for each vertex w in v's neighbors
        if (w has not been visited)
            dfTrav(w, v)
```

- Java method:

```
private static void dfTrav(Vertex v, Vertex parent) {
    System.out.println(v.id);    // visit v
    v.done = true;
    v.parent = parent;
    Edge e = v.edges;
    while (e != null) {
        Vertex w = e.end;
        if (!w.done)
            dfTrav(w, v);
        e = e.next;
    }
}
```

## Example: Depth-First Traversal from Portland





For the examples, we'll assume that the edges in each vertex's adjacency list are sorted by increasing edge cost.

```
void dfTrav(Vertex v, Vertex parent) {
    System.out.println(v.id);
    v.done = true;
    v.parent = parent;
    Edge e = v.edges;
    while (e != null) {
        Vertex w = e.end;
        if (!w.done)
            dfTrav(w, v);
        e = e.next;
    }
}
```

```
dfTrav(Ptl, null)
w = Pts
dfTrav(Pts, Ptl)
w = Ptl, Bos
dfTrav(Bos, Pts)
w = Wor
dfTrav(Wor, Bos)
w = Pro
dfTrav(Pro, Wor)
w = Wor, Bos, NY
dfTrav(NY, Pro)
w = Pro
return
no more neighbors
return
w = Bos, Con
dfTrav(Con, Wor)
...
```

## Depth-First Spanning Tree




The edges obtained by following the parent references form a spanning tree with the origin of the traversal as its root.

From any city, we can get to the origin by following the roads in the spanning tree.




## Another Example: Depth-First Traversal from Worcester


- In what order will the cities be visited?
- Which edges will be in the resulting spanning tree?



## Checking for Cycles in an Undirected Graph



- To discover a cycle in an undirected graph, we can:
  - perform a depth-first traversal, marking the vertices as visited
  - when considering neighbors of a visited vertex, if we discover one already marked as visited, there must be a cycle
- If no cycles found during the traversal, the graph is acyclic.
- This doesn't work for directed graphs:
  - c is a neighbor of both a and b
  - there is no cycle



## Breadth-First Traversal


- Use a queue, as we did for BFS and level-order tree traversal:

```
private static void bfTrav(Vertex origin) {
    origin.encountered = true;
    origin.parent = null;
    Queue<Vertex> q = new LLQueue<Vertex>();
    q.insert(origin);

    while (!q.isEmpty()) {
        Vertex v = q.remove();
        System.out.println(v.id);           // Visit v.

        // Add v's unencountered neighbors to the queue.
        Edge e = v.edges;
        while (e != null) {
            Vertex w = e.end;
            if (!w.encountered) {
                w.encountered = true;
                w.parent = v;
                q.insert(w);
            }
            e = e.next;
        }
    }
}
```


### Example: Breadth-First Traversal from Portland




Evolution of the queue:

<u>remove</u>	<u>insert</u>	<u>queue contents</u>
	Portland	Portland
Portland	Portsmouth, Concord	Portsmouth, Concord
Portsmouth	Boston, Worcester	Concord, Boston, Worcester
Concord	none	Boston, Worcester
Boston	Providence	Worcester, Providence
Worcester	Albany	Providence, Albany
Providence	New York	Albany, New York
Albany	none	New York
New York	none	empty


## Breadth-First Spanning Tree




**breadth-first spanning tree:**



**depth-first spanning tree:**



## Another Example: Breadth-First Traversal from Worcester



Evolution of the queue:

remove

insert

queue contents

## Time Complexity of Graph Traversals

- let  $V$  = number of vertices in the graph  
 $E$  = number of edges
- If we use an adjacency matrix, a traversal requires  $O(V^2)$  steps.
  - why?
- If we use an adjacency list, a traversal requires  $O(V + E)$  steps.
  - visit each vertex once
  - traverse each vertex's adjacency list at most once
    - the total length of the adjacency lists is at most  $2E = O(E)$
  - $O(V + E) \ll O(V^2)$  for a sparse graph
  - for a dense graph,  $E = O(V^2)$ , so both representations are  $O(V^2)$
- In our implementations of the remaining algorithms, we'll assume an adjacency-list implementation.

## Minimum Spanning Tree

- A minimum spanning tree (MST) has the smallest total cost among all possible spanning trees.
  - *example:*


one possible spanning tree  
(total cost =  $39 + 83 + 54 = 176$ )

the minimal-cost spanning tree  
(total cost =  $39 + 54 + 44 = 137$ )
- If no two edges have the same cost, there is a unique MST.  
If two or more edges have the same cost, there may be more than one MST.
- Finding an MST could be used to:
  - determine the shortest highway system for a set of cities
  - calculate the smallest length of cable needed to connect a network of computers

## Building a Minimum Spanning Tree

- Key insight: if you divide the vertices into two disjoint subsets A and B, then the lowest-cost edge joining a vertex in A to a vertex in B – call it (a, b) – must be part of the MST.


• *example:*



The 6 bold edges each join an unshaded vertex to a shaded vertex.

The one with the lowest cost (Portland to Portsmouth) must be in the MST.


- Proof by contradiction:
  - assume there is an MST (call it T) that doesn't include (a, b)
  - T must include a path from a to b, so it must include one of the other edges (a', b') that spans subsets A and B, such that (a', b') is part of the path from a to b
  - adding (a, b) to T introduces a cycle
  - removing (a', b') gives a spanning tree with lower cost, which contradicts the original assumption.



## Prim's MST Algorithm

- Begin with the following subsets:
  - A = any one of the vertices
  - B = all of the other vertices
- Repeatedly select the lowest-cost edge (a, b) connecting a vertex in A to a vertex in B and do the following:
  - add (a, b) to the spanning tree
  - update the two sets:  $A = A \cup \{b\}$   
 $B = B - \{b\}$
- Continue until A contains all of the vertices.


## Example: Prim's Starting from Concord



- Tracing the algorithm:

edge added	set A	set B
	{Con}	{Alb, Bos, NY, Ptl, Pts, Pro, Wor}
(Con, Wor)	{Con, Wor}	{Alb, Bos, NY, Ptl, Pts, Pro}
(Wor, Pro)	{Con, Wor, Pro}	{Alb, Bos, NY, Ptl, Pts}
(Wor, Bos)	{Con, Wor, Pro, Bos}	{Alb, NY, Ptl, Pts}
(Bos, Pts)	{Con, Wor, Pro, Bos, Pts}	{Alb, NY, Ptl}
(Pts, Ptl)	{Con, Wor, Pro, Bos, Pts, Ptl}	{Alb, NY}
(Wor, Alb)	{Con, Wor, Pro, Bos, Pts, Ptl, Alb}	{NY}
(Pro, NY)	{Con, Wor, Pro, Bos, Pts, Ptl, Alb, NY}	{}

## MST May Not Give Shortest Paths



- The MST is the spanning tree with the minimal *total* edge cost.
- It does not necessarily include the minimal cost path between a pair of vertices.
- Example: shortest path from Boston to Providence is along the single edge connecting them
  - that edge is not in the MST

## Implementing Prim's Algorithm in our Graph class

- Use the done field to keep track of the sets.
  - if  $v.done == true$ ,  $v$  is in set A
  - if  $v.done == false$ ,  $v$  is in set B
- Repeatedly scan through the lists of vertices and edges to find the next edge to add.
  - $O(EV)$
- We can do better!
  - use a heap-based priority queue to store the vertices in set B
  - priority of a vertex  $x = -1 * \text{cost of the lowest-cost edge connecting } x \text{ to a vertex in set A}$ 
    - why multiply by  $-1$ ?
  - somewhat tricky: need to update the priorities over time
    - $O(E \log V)$


## The Shortest-Path Problem

- It's often useful to know the shortest path from one vertex to another – i.e., the one with the minimal total cost
  - example application: routing traffic in the Internet
- For an *unweighted* graph, we can simply do the following:
  - start a breadth-first traversal from the origin,  $v$
  - stop the traversal when you reach the other vertex,  $w$
  - the path from  $v$  to  $w$  in the resulting (possibly partial) spanning tree is a shortest path
- A breadth-first traversal works for an unweighted graph because:
  - the shortest path is simply one with the fewest edges
  - a breadth-first traversal visits cities in order according to the number of edges they are from the origin.
- Why might this approach fail to work for a *weighted* graph?

## Dijkstra's Algorithm


- One algorithm for solving the shortest-path problem for weighted graphs was developed by E.W. Dijkstra.
- It allows us to find the shortest path from a vertex  $v$  (the origin) to *all other vertices* that can be reached from  $v$ .
- Basic idea:
  - maintain estimates of the shortest paths from the origin to every vertex (along with their costs)
  - gradually refine these estimates as we traverse the graph
- Initial estimates:

	<u>path</u>	<u>cost</u>
the origin itself:	stay put!	0
all other vertices:	unknown	infinity



## Dijkstra's Algorithm (cont.)

- We say that a vertex  $w$  is *finalized* if we have found the shortest path from  $v$  to  $w$ .
- We repeatedly do the following:
  - find the unfinalized vertex  $w$  with the lowest cost estimate
  - mark  $w$  as finalized (shown as a filled circle below)
  - examine each unfinalized neighbor  $x$  of  $w$  to see if there is a shorter path to  $x$  that passes through  $w$ 
    - if there is, update the shortest-path estimate for  $x$
- Example:






## Another Example: Shortest Paths from Providence

- Initial estimates:


Boston	infinity
Worcester	infinity
Portsmouth	infinity
Providence	<b>0</b>



- Providence has the smallest unfinalized estimate, so we finalize it.


- We update our estimates for its neighbors:

Boston	49 (< infinity)
Worcester	42 (< infinity)
Portsmouth	infinity
Providence	<b>0</b>



## Shortest Paths from Providence (cont.)


Boston	49
Worcester	<b>42</b>
Portsmouth	infinity
Providence	<b>0</b>



- Worcester has the smallest unfinalized estimate, so we finalize it.
  - any other route from Prov. to Worc. would need to go via Boston, and since  $(\text{Prov} \rightarrow \text{Worc}) < (\text{Prov} \rightarrow \text{Bos})$ , we can't do better.


- We update our estimates for Worcester's unfinalized neighbors:

Boston	49 (no change)
Worcester	<b>42</b>
Portsmouth	125 ( $42 + 83 < \text{infinity}$ )
Providence	<b>0</b>



### Shortest Paths from Providence (cont.)


Boston	49
Worcester	42
Portsmouth	125
Providence	0



- Boston has the smallest unfinalized estimate, so we finalize it.
  - we'll see later why we can safely do this!


- We update our estimates for Boston's unfinalized neighbors:

Boston	49
Worcester	42
Portsmouth	103 (49 + 54 < 125)
Providence	0




### Shortest Paths from Providence (cont.)

Boston	49
Worcester	42
Portsmouth	103
Providence	0



- Only Portsmouth is left, so we finalize it.

## Finalizing a Vertex



- Let  $w$  be the unfinalized vertex with the smallest cost estimate. Why can we finalize  $w$ , before seeing the rest of the graph?
- We know that  $w$ 's current estimate is for the shortest path to  $w$  that passes through only *finalized* vertices.
- Any shorter path to  $w$  would have to pass through one of the other encountered-but-unfinalized vertices, but we know that they're all further away from the origin than  $w$  is.
  - their cost estimates may decrease in subsequent stages of the algorithm, but they can't drop below  $w$ 's current estimate!

## Pseudocode for Dijkstra's Algorithm


- ```

dijkstra(origin)
  origin.cost = 0
  for each other vertex  $v$ 
     $v$ .cost = infinity;

  while there are still unfinalized vertices with cost < infinity
    find the unfinalized vertex  $w$  with the minimal cost
    mark  $w$  as finalized

    for each unfinalized vertex  $x$  adjacent to  $w$ 
       $\text{cost\_via\_w} = w.\text{cost} + \text{edge\_cost}(w, x)$ 
      if ( $\text{cost\_via\_w} < x.\text{cost}$ )
         $x.\text{cost} = \text{cost\_via\_w}$ 
         $x.\text{parent} = w$ 
  
```
- At the conclusion of the algorithm, for each vertex  $v$ :
    - $v.\text{cost}$  is the cost of the shortest path from the origin to  $v$ ;
    - if  $v.\text{cost}$  is infinity, there is no path from the origin to  $v$
    - starting at  $v$  and following the parent references yields the shortest path

### Example: Shortest Paths from Concord




Evolution of the cost estimates (costs in bold have been finalized):

|            |          |           |           |           |            |            |            |            |
|------------|----------|-----------|-----------|-----------|------------|------------|------------|------------|
| Albany     | inf      | inf       | 197       | 197       | 197        | 197        | <b>197</b> |            |
| Boston     | inf      | 74        | <b>74</b> |           |            |            |            |            |
| Concord    | <b>0</b> |           |           |           |            |            |            |            |
| New York   | inf      | inf       | inf       | inf       | inf        | 290        | 290        | <b>290</b> |
| Portland   | inf      | 84        | 84        | <b>84</b> |            |            |            |            |
| Portsmouth | inf      | inf       | 146       | 128       | 123        | <b>123</b> |            |            |
| Providence | inf      | inf       | 105       | 105       | <b>105</b> |            |            |            |
| Worcester  | inf      | <b>63</b> |           |           |            |            |            |            |

*Note that the Portsmouth estimate was improved three times!*

### Another Example: Shortest Paths from Worcester



Evolution of the cost estimates (costs in bold have been finalized):


|            |  |  |  |  |  |  |  |  |
|------------|--|--|--|--|--|--|--|--|
| Albany     |  |  |  |  |  |  |  |  |
| Boston     |  |  |  |  |  |  |  |  |
| Concord    |  |  |  |  |  |  |  |  |
| New York   |  |  |  |  |  |  |  |  |
| Portland   |  |  |  |  |  |  |  |  |
| Portsmouth |  |  |  |  |  |  |  |  |
| Providence |  |  |  |  |  |  |  |  |
| Worcester  |  |  |  |  |  |  |  |  |

## Implementing Dijkstra's Algorithm

- Similar to the implementation of Prim's algorithm.
- Use a heap-based priority queue to store the unfinalized vertices.
  - priority = ?
- Need to update a vertex's priority whenever we update its shortest-path estimate.
- Time complexity =  $O(E \log V)$

## Topological Sort

- Used to order the vertices in a directed acyclic graph (a DAG).
- Topological order: an ordering of the vertices such that, if there is directed edge from a to b, a comes before b.
- Example application: ordering courses according to prerequisites



- a directed edge from a to b indicates that a is a prereq of b
- There may be more than one topological ordering.


## Topological Sort Algorithm

- A *successor* of a vertex  $v$  in a directed graph = a vertex  $w$  such that  $(v, w)$  is an edge in the graph ( $v \rightarrow w$ )
- Basic idea: find vertices that have no successors and work backward from them.
  - there must be at least one such vertex. why?
- Pseudocode for one possible approach:
 

```

topoSort
    S = a stack to hold the vertices as they are visited
    while there are still unvisited vertices
        find a vertex v with no unvisited successors
        mark v as visited
        S.push(v)
    return S
            
```
- Popping the vertices off the resulting stack gives one possible topological ordering.

## Topological Sort Example




Evolution of the stack:

| push  | stack contents (top to bottom)                                      |
|-------|---------------------------------------------------------------------|
| E-124 | E-124                                                               |
| E-162 | E-162, E-124                                                        |
| E-215 | E-215, E-162, E-124                                                 |
| E-104 | E-104, E-215, E-162, E-124                                          |
| E-119 | E-119, E-104, E-215, E-162, E-124                                   |
| E-160 | E-160, E-119, E-104, E-215, E-162, E-124                            |
| E-10  | E-10, E-160, E-119, E-104, E-215, E-162, E-124                      |
| E-50b | E-50b, E-10, E-160, E-119, E-104, E-215, E-162, E-124               |
| E-50a | <b>E-50a, E-50b, E-10, E-160, E-119, E-104, E-215, E-162, E-124</b> |

one possible topological ordering

## Another Topological Sort Example




Evolution of the stack:

push


stack contents (top to bottom)

## Traveling Salesperson Problem (TSP)



- A salesperson needs to travel to a number of cities to visit clients, and wants to do so as efficiently as possible.
- As in our earlier problems, we use a weighted graph.
- A *tour* is a path that begins at some starting vertex, passes through every other vertex *once and only once*, and returns to the starting vertex. (The actual starting vertex doesn't matter.)
- TSP: find the tour with the lowest total cost
- TSP algorithms assume the graph is complete, but we can assign infinite costs if there isn't a direct route between two cities.

## TSP for Santa Claus



source: <http://www.tsp.gatech.edu/world/pictures.html>


A "world TSP" with 1,904,711 cities.

The figure at right shows a tour with a total cost of 7,516,353,779 meters – which is at most 0.068% longer than the optimal tour.

- Other applications:
  - coin collection from phone booths
  - routes for school buses or garbage trucks
  - minimizing the movements of machines in automated manufacturing processes
  - many others

## Solving a TSP: Brute-Force Approach

- Perform an exhaustive search of all possible tours.
- One way: use DFS to traverse the entire state-space search tree.



- The leaf nodes correspond to possible solutions.
  - for  $n$  cities, there are  $(n - 1)!$  leaf nodes in the tree.
  - half are redundant (e.g., L-Cm-Ct-O-Y-L = L-Y-O-Ct-Cm-L)
- Problem: exhaustive search is intractable for all but small  $n$ .
  - example: when  $n = 14$ ,  $((n - 1)!)/2 =$  over 3 billion



## Solving a TSP: Informed State-Space Search

- Use A\* with an appropriate heuristic function for estimating the cost of the remaining edges in the tour.
- This is much better than brute force, but it still uses exponential space and time.

## Algorithm Analysis Revisited

- Recall that we can group algorithms into classes ( $n$  = problem size):

| <u>name</u>       | <u>example expressions</u>     | <u>big-O notation</u> |
|-------------------|--------------------------------|-----------------------|
| constant time     | 1, 7, 10                       | $O(1)$                |
| logarithmic time  | $3 \lg_{10} n$ , $\lg_2 n + 5$ | $O(\lg n)$            |
| linear time       | $5n$ , $10n - 2 \lg_2 n$       | $O(n)$                |
| $n \log n$ time   | $4n \lg_2 n$ , $n \lg_2 n + n$ | $O(n \lg n)$          |
| quadratic time    | $2n^2 + 3n$ , $n^2 - 1$        | $O(n^2)$              |
| $n^c$ ( $c > 2$ ) | $n^3 - 5n$ , $2n^5 + 5n^2$     | $O(n^c)$              |
| exponential time  | $2^n$ , $5e^n + 2n^2$          | $O(c^n)$              |
| factorial time    | $(n-1)! / 2$ , $3n!$           | $O(n!)$               |

- Algorithms that fall into one of the classes above the dotted line are referred to as *polynomial-time* algorithms.
- The term *exponential-time algorithm* is sometimes used to include *all* algorithms that fall below the dotted line.
  - algorithms whose running time grows as fast or faster than  $c^n$

## Classifying Problems

- Problems that can be solved using a polynomial-time algorithm are considered “easy” problems.
  - we can solve large problem instances in a reasonable amount of time
- Problems that don't have a polynomial-time solution algorithm are considered “hard” or “intractable” problems.
  - they can only be solved exactly for small values of  $n$
- Increasing the CPU speed doesn't help much for intractable problems:

|                                  | CPU 1 | CPU 2<br>(1000x faster) |
|----------------------------------|-------|-------------------------|
| max problem size for $O(n)$ alg: | $N$   | $1000N$                 |
| $O(n^2)$ alg:                    | $N$   | $31.6 N$                |
| $O(2^n)$ alg:                    | $N$   | $N + 9.97$              |

## Classifying Problems (cont.)

- The class of problems that can be solved using a polynomial-time algorithm is known as the class P.
- Many problems that don't have a polynomial-time solution algorithm belong to a class known as NP
  - for *non-deterministic polynomial*
- If a problem is in NP, it's possible to guess a solution and verify if the guess is correct in polynomial time.
  - example: a variant of the TSP in which we attempt to determine if there is a tour with total cost  $\leq$  some bound  $b$
  - given a tour, it takes polynomial time to add up the costs of the edges and compare the result to  $b$

### Classifying Problems (cont.)

- If a problem is *NP-complete*, then finding a polynomial-time solution for it would allow you to solve a large number of other hard problems.
  - thus, it's extremely unlikely such a solution exists!
- The TSP variant described on the previous slide is NP-complete.
- Finding the optimal tour is at least as hard.
- For more info. about problem classes, there is a good video of a lecture by Prof. Michael Sipser of MIT available here:  
<http://claymath.msri.org/sipser2006.mov>

### Dealing With Intractable Problems

- When faced with an intractable problem, we resort to techniques that quickly find solutions that are "good enough".
- Such techniques are often referred to as *heuristic* techniques.
  - heuristic = rule of thumb
  - there's no guarantee these techniques will produce the optimal solution, but they typically work well

## Iterative Improvement Algorithms

- One type of heuristic technique is what is known as an *iterative improvement algorithm*.
  - start with a randomly chosen solution
  - gradually make small changes to the solution in an attempt to improve it
    - e.g., change the position of one label
  - stop after some number of iterations
- There are several variations of this type of algorithm.

## Hill Climbing

- Hill climbing is one type of iterative improvement algorithm.
  - start with a randomly chosen solution
  - repeatedly consider possible small changes to the solution
  - if a change would improve the solution, make it
  - if a change would make the solution worse, don't make it
  - stop after some number of iterations
- It's called hill climbing because it repeatedly takes small steps that improve the quality of the solution.
  - "climbs" towards the optimal solution

## Simulated Annealing

- *Simulated annealing* is another iterative improvement algorithm.
  - start with a randomly chosen solution
  - repeatedly consider possible small changes to the solution
  - if a change would improve the solution, make it
  - if a change would make the solution worse, *make it some of the time* (according to some probability)
  - the probability of doing so reduces over time

## Take-Home Lessons

- Object-oriented programming allows us to capture the abstractions in the programs that we write.
  - creates reusable building blocks
  - key concepts: encapsulation, inheritance, polymorphism
- Abstract data types allow us to organize and manipulate collections of data.
  - a given ADT can be implemented in different ways
  - fundamental building blocks: arrays, linked nodes
- Efficiency matters when dealing with large collections of data.
  - some solutions can be *much* faster or more space efficient than others!
  - what's the best data structure/algorithm for the specific instances of the problem that you expect to see?
    - example: sorting an almost sorted collection

## Take-Home Lessons (cont.)

- Use the tools in your toolbox!
  - generic data structures
  - lists/stacks/queues
  - trees
  - heaps
  - hash tables
  - recursion
  - recursive backtracking
  - divide-and-conquer
  - state-space search
  - ...

## From the Introductory Lecture...

- We will study fundamental *data structures*.
  - ways of imposing order on a collection of information
  - sequences: lists, stacks, and queues
  - trees
  - hash tables
  - graphs
- We will also:
  - study *algorithms* related to these data structures
  - learn how to *compare* data structures & algorithms
- Goals:
  - learn to think more intelligently about programming problems
  - acquire a set of useful tools and techniques