

# Heaps and Priority Queues

Computer Science E-22  
Harvard Extension School


David G. Sullivan, Ph.D.

## Priority Queue


- A *priority queue* is a collection in which each item in the collection has an associated number known as a *priority*.
  - ("Henry Leitner", 10), ("Drew Faust", 15), ("Dave Sullivan", 5)
  - use a higher priority for items that are "more important"
- Example: scheduling a shared resource like the CPU
  - give some processes/applications a higher priority, so that they will be scheduled first and/or more often
- Key operations:
  - *insert*: add an item to the priority queue, positioning it according to its priority
  - *remove*: remove the item with the highest priority
- How can we efficiently implement a priority queue?
  - use a type of binary tree known as a *heap*

## Complete Binary Trees

- A binary tree of height  $h$  is *complete* if:
  - levels 0 through  $h - 1$  are fully occupied
  - there are no “gaps” to the left of a node in level  $h$
- Complete:




- Not complete (○ = missing node):





## Representing a Complete Binary Tree

- A complete binary tree has a simple array representation.
- The nodes of the tree are stored in the array in the order in which they would be visited by a level-order traversal (i.e., top to bottom, left to right).




- Examples:



## Navigating a Complete Binary Tree in Array Form


- The root node is in  $a[0]$
- Given the node in  $a[i]$ :
  - its left child is in  $a[2*i + 1]$
  - its right child is in  $a[2*i + 2]$
  - its parent is in  $a[(i - 1)/2]$  (using integer division)



- Examples:
  - the left child of the node in  $a[1]$  is in  $a[2*1 + 1] = a[3]$
  - the right child of the node in  $a[3]$  is in  $a[2*3 + 2] = a[8]$
  - the parent of the node in  $a[4]$  is in  $a[(4-1)/2] = a[1]$
  - the parent of the node in  $a[7]$  is in  $a[(7-1)/2] = a[3]$

## Heaps

- Heap: a complete binary tree in which each interior node is greater than or equal to its children
- Examples:



- The largest value is always at the root of the tree.
- The smallest value can be in *any* leaf node – there's no guarantee about which one it will be.
- Strictly speaking, the heaps that we will use are *max-at-top* heaps. You can also define a *min-at-top* heap, in which every interior node is less than or equal to its children.

## How to Compare Objects

- We need to be able to compare items in the heap.
- If those items are objects, we can't just do something like this:

```
if (item1 < item2)
```

Why not?

- Instead, we need to use a method to compare them.

## An Interface for Objects That Can Be Compared

- The Comparable interface is a built-in generic Java interface:

```
public interface Comparable<T> {  
    public int compareTo(T other);  
}
```

- It is used when defining a class of objects that can be ordered.
- Examples from the built-in Java classes:

```
public class String implements Comparable<String> {  
    ...  
    public int compareTo(String other) {  
        ...  
    }  
    public class Integer implements Comparable<Integer> {  
        ...  
        public int compareTo(Integer other) {  
            ...  
        }  
    }
```

## An Interface for Objects That Can Be Compared (cont.)

```
public interface Comparable<T> {  
    public int compareTo(T other);  
}
```

- `item1.compareTo(item2)` should return:
  - a negative integer if `item1` "comes before" `item2`
  - a positive integer if `item1` "comes after" `item2`
  - 0 if `item1` and `item2` are equivalent in the ordering
- These conventions make it easy to construct appropriate method calls:

### numeric comparison

`item1 < item2`

`item1 > item2`

`item1 == item2`

### comparison using compareTo

`item1.compareTo(item2) < 0`

`item1.compareTo(item2) > 0`

`item1.compareTo(item2) == 0`

## A Class for Items in a Priority Queue


```
public class PQItem implements Comparable<PQItem> {  
    // group an arbitrary object with a priority  
    private Object data;  
    private int priority;  
    ...  
  
    public int compareTo(PQItem other) {  
        // error-checking goes here...  
        return (priority - other.priority);  
    }  
}
```

- Its `compareTo()` compares `PQItem`s based on their priorities.
- `item1.compareTo(item2)` returns:
  - a negative integer if `item1` has a lower priority than `item2`
  - a positive integer if `item1` has a higher priority than `item2`
  - 0 if they have the same priority

## Heap Implementation

```
public class Heap<T extends Comparable<T>> {
    private T[] contents;
    private int numItems;

    public Heap(int maxSize) {
        contents = (T[])new Comparable[maxSize];
        numItems = 0;
    }
    ...
}
```




- Heap is another example of a generic collection class.
  - as usual, T is the type of the elements
  - extends `Comparable<T>` specifies T must implement `Comparable<T>`
  - must use `Comparable` (not `Object`) when creating the array

## Heap Implementation (cont.)

```
public class Heap<T extends Comparable<T>> {
    private T[] contents;
    private int numItems;

    ...
}
```



- The picture above is a heap of integers:
 

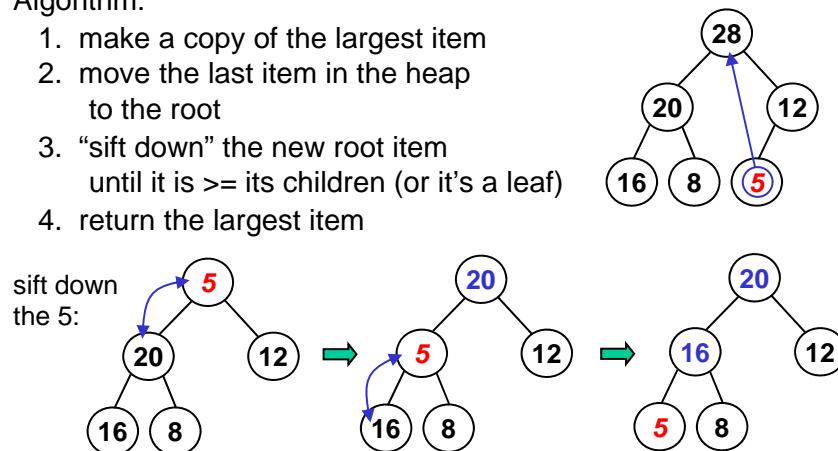
```
Heap<Integer> myHeap = new Heap<Integer>(20);
```

  - works because `Integer` implements `Comparable<Integer>`
  - could also use `String` or `Double`
- For a priority queue, we can use objects of our `PQItem` class:
 

```
Heap<PQItem> pqueue = new Heap<PQItem>(50);
```

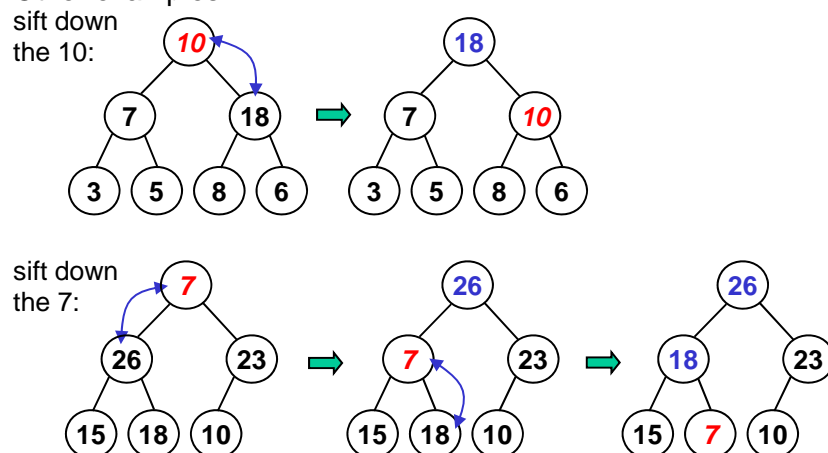
## Removing the Largest Item from a Heap

- Remove and return the item in the root node.
- In addition, we need to move the largest remaining item to the root, while maintaining a complete tree with each node  $\geq$  children
- Algorithm:
  - make a copy of the largest item
  - move the last item in the heap to the root
  - "sift down" the new root item until it is  $\geq$  its children (or it's a leaf)
  - return the largest item



## Sifting Down an Item

- To sift down item  $x$  (i.e., the item whose key is  $x$ ):
  - compare  $x$  with the larger of the item's children,  $y$
  - if  $x < y$ , swap  $x$  and  $y$  and repeat
- Other examples:




## siftDown() Method

```
private void siftDown(int i) {
    T toSift = contents[i];
    int parent = i;
    int child = 2 * parent + 1;
    while (child < numItems) {
        // If the right child is bigger, compare with it.
        if (child < numItems - 1 &&
            contents[child].compareTo(contents[child + 1]) < 0)
            child = child + 1;

        if (toSift.compareTo(contents[child]) >= 0)
            break; // we're done

        // Move child up and move down one level in the tree.
        contents[parent] = contents[child];
        parent = child;
        child = 2 * parent + 1;
    }
    contents[parent] = toSift;
}
```

- We don't actually swap items. We wait until the end to put the sifted item in place.



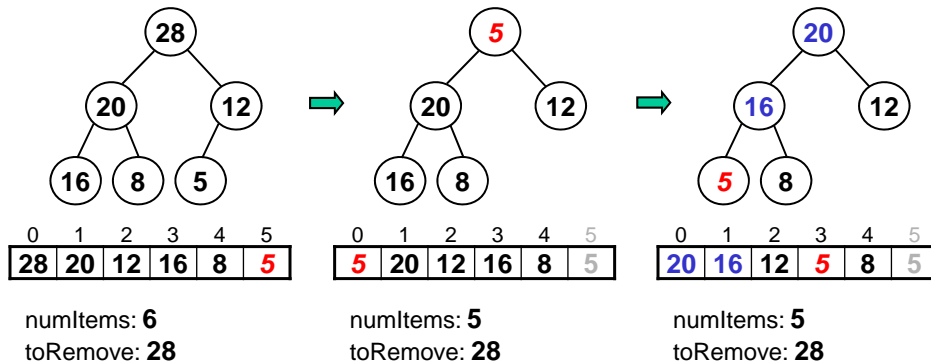
toSift: 7	
parent	child
0	1
1	3
1	4
4	9

## remove() Method

```
public T remove() {
    T toRemove = contents[0];

    contents[0] = contents[numItems - 1];
    numItems--;
    siftDown(0);

    return toRemove;
}
```






## Inserting an Item in a Heap


- Algorithm:
  1. put the item in the next available slot (grow array if needed)
  2. "sift up" the new item until it is  $\leq$  its parent (or it becomes the root item)

- Example: insert 35

put it in place:

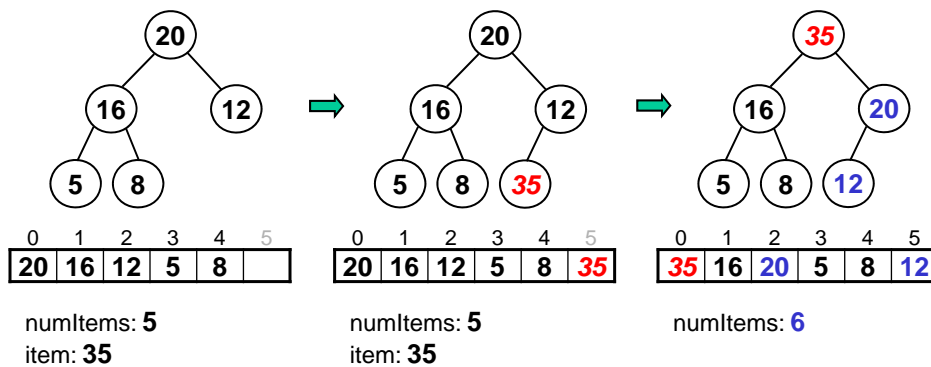


sift it up:



## insert() Method

```
public void insert(T item) {
    if (numItems == contents.length) {
        // code to grow the array goes here...
    }
    contents[numItems] = item;
    siftUp(numItems);
    numItems++;
}
```




## Converting an Arbitrary Array to a Heap


- Algorithm to convert an array with  $n$  items to a heap:
  - start with the parent of the last element:  
 $\text{contents}[i]$ , where  $i = ((n - 1) - 1) / 2 = (n - 2) / 2$
  - sift down  $\text{contents}[i]$  and all elements to its left

- Example:

0	1	2	3	4	5	6
5	16	8	14	20	1	26




- Last element's parent =  $\text{contents}[(7 - 2) / 2] = \text{contents}[2]$ .  
Sift it down:




## Converting an Array to a Heap (cont.)

- Next, sift down  $\text{contents}[1]$ :



- Finally, sift down  $\text{contents}[0]$ :




## Creating a Heap from an Array

```
public class Heap<T extends Comparable<T>> {
    private T[] contents;
    private int numItems;
    ...

    public Heap(T[] arr) {
        // Note that we don't copy the array!
        contents = arr;
        numItems = arr.length;
        makeHeap();
    }

    private void makeHeap() {
        int last = contents.length - 1;
        int parentOfLast = (last - 1)/2;
        for (int i = parentOfLast; i >= 0; i--)
            siftDown(i);
        ...
    }
}
```

## Time Complexity of a Heap



- A heap containing  $n$  items has a height  $\leq \log_2 n$ .
- Thus, removal and insertion are both  $O(\log n)$ .
  - remove: go down at most  $\log_2 n$  levels when sifting down from the root, and do a constant number of operations per level
  - insert: go up at most  $\log_2 n$  levels when sifting up to the root, and do a constant number of operations per level
- This means we can use a heap for a  $O(\log n)$ -time priority queue.
- Time complexity of creating a heap from an array?

## Using a Heap to Sort an Array

- Recall selection sort: it repeatedly finds the smallest remaining element and swaps it into place:

0	1	2	3	4	5	6
5	16	8	14	20	1	26
0	1	2	3	4	5	6
1	16	8	14	20	5	26
0	1	2	3	4	5	6
1	5	8	14	20	16	26

...

- It isn't efficient ( $O(n^2)$ ), because it performs a linear scan to find the smallest remaining element ( $O(n)$  steps per scan).
- Heapsort is a sorting algorithm that repeatedly finds the *largest* remaining element and puts it in place.
- It *is* efficient ( $O(n \log n)$ ), because it turns the array into a heap, which means that it can find and remove the largest remaining element in  $O(\log n)$  steps.

## Heapsort


```
public class HeapSort {
    public static <T extends Comparable<T>> void
    heapSort(T[] arr) {
        // Turn the array into a max-at-top heap.
        Heap<T> heap = new Heap<T>(arr);
        int endUnsorted = arr.length - 1;
        while (endUnsorted > 0) {
            // Get the largest remaining element and put it
            // at the end of the unsorted portion of the array.
            T largestRemaining = heap.remove();
            arr[endUnsorted] = largestRemaining;
            endUnsorted--;
        }
    }
}
```

- We define a *generic method*, with a type variable in the method header. It goes right before the method's return type.
- T is a placeholder for the type of the array.
  - can be any type T that implements Comparable<T>.


## Heapsort Example

- Sort the following array: 

0	1	2	3	4	5	6
13	6	45	10	3	22	5
- Here's the corresponding complete tree:




- Begin by converting it to a heap:



## Heapsort Example (cont.)

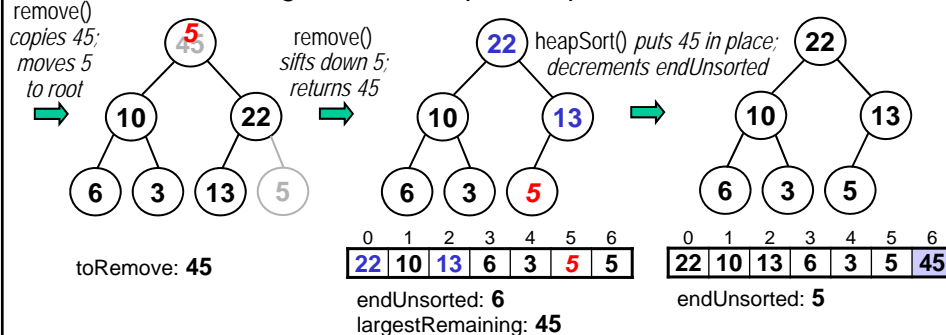
- Here's the heap in both tree and array forms:



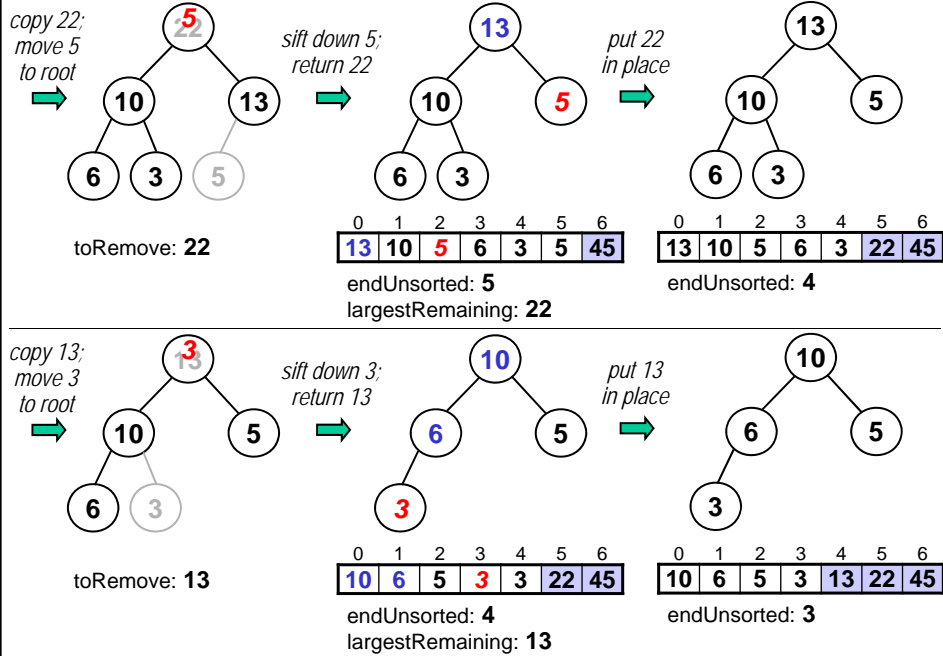
0	1	2	3	4	5	6
45	10	22	6	3	13	5

endUnsorted: 6

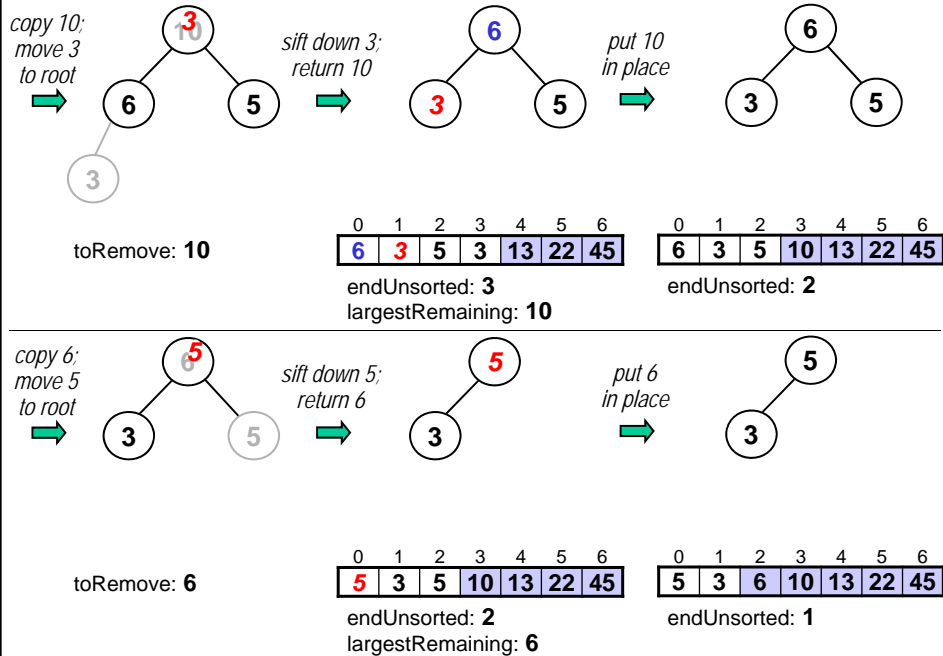
- Remove the largest item and put it in place:




## Heapsort Example (cont.)



## Heapsort Example (cont.)






### How Does Heapsort Compare?

algorithm	best case	avg case	worst case	extra memory
selection sort	$O(n^2)$	$O(n^2)$	$O(n^2)$	$O(1)$
insertion sort	$O(n)$	$O(n^2)$	$O(n^2)$	$O(1)$
Shell sort	$O(n \log n)$	$O(n^{1.5})$	$O(n^{1.5})$	$O(1)$
bubble sort	$O(n^2)$	$O(n^2)$	$O(n^2)$	$O(1)$
quicksort	$O(n \log n)$	$O(n \log n)$	$O(n^2)$	$O(1)$
mergesort	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$	$O(n)$
<b>heapsort</b>	<b><math>O(n \log n)</math></b>	<b><math>O(n \log n)</math></b>	<b><math>O(n \log n)</math></b>	<b><math>O(1)</math></b>

- Heapsort matches mergesort for the best worst-case time complexity, but it has better space complexity.
- Insertion sort is still best for arrays that are almost sorted.
  - heapsort will scramble an almost sorted array before sorting it
- Quicksort is still typically fastest in the average case.

## State-Space Search Revisited

- Earlier, we considered three algorithms for state-space search:
  - breadth-first search (BFS)
  - depth-first search (DFS)
  - iterative-deepening search (IDS)
- These are all *uninformed* search algorithms.
  - always consider the states in a certain order
  - do not consider how close a given state is to the goal
- 8 Puzzle example:



## Informed State-Space Search

- *Informed* search algorithms attempt to consider more promising states first.
- These algorithms associate a *priority* with each successor state that is generated.
  - base priority on an estimate of nearness to a goal state
  - when choosing the next state to consider, select the one with the highest priority
- Use a priority queue to store the yet-to-be-considered search nodes.




## State-Space Search: Estimating the Remaining Cost

- The priority of a state is based on the *remaining cost* – i.e., the cost of getting from the state to the closest goal state.
  - for the 8 puzzle, remaining cost = # of steps to closest goal
- For most problems, we can't determine the exact remaining cost.
  - if we could, we wouldn't need to search!
- Instead, we estimate the remaining cost using a *heuristic function*  $h(x)$  that takes a state  $x$  and computes a cost estimate for it.
  - heuristic = rule of thumb
- To find optimal solutions, we need an *admissible* heuristic – one that never overestimates the remaining cost.


## Heuristic Function for the Eight Puzzle

- Manhattan distance = horizontal distance + vertical distance

- example: For the board at right, the Manhattan distance of the **3** tile from its position in the goal state = 1 column + 1 row = 2



- Use  $h(x)$  = sum of the Manhattan distances of the tiles in  $x$  from their positions in the goal state
  - for our example:




$$h(x) = 1 + 1 + 2 + 2 + 1 + 0 + 1 + 1 = 9$$

- This heuristic is admissible because each of the operators (move blank up, move blank down, etc.) moves a single tile a distance of 1, so it will take at least  $h(x)$  steps to reach the goal.

## Greedy Search

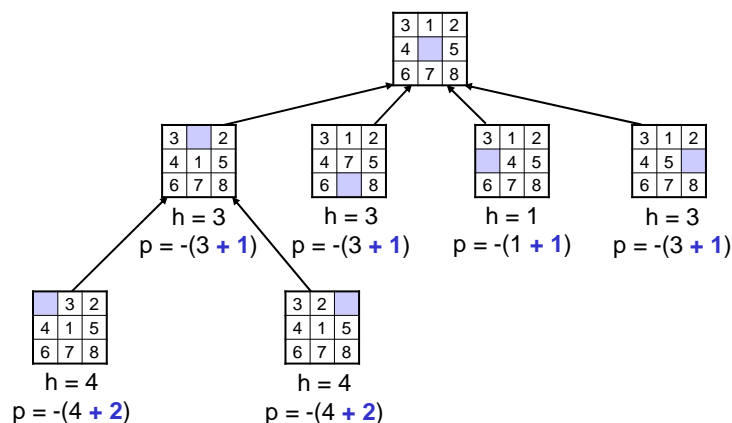
- Priority of state  $x$ ,  $p(x) = -1 * h(x)$ 
  - mult. by  $-1$  so states closer to the goal have higher priorities



- Greedy search would consider the highlighted successor before the other successors, because it has the highest priority.
- Greedy search is:
  - incomplete: it may not find a solution
    - it could end up going down an infinite path
  - not optimal: the solution it finds may not have the lowest cost
    - it fails to consider the cost of getting to the current state

## A\* Search

- Priority of state  $x$ ,  $p(x) = -1 * (h(x) + g(x))$   
where  $g(x)$  = the cost of getting from the initial state to  $x$



- Incorporating  $g(x)$  allows A\* to find an optimal solution – one with the minimal *total* cost.

## Characteristics of A\*

- It is complete and optimal.
  - provided that  $h(x)$  is admissible, and that  $g(x)$  increases or stays the same as the depth increases
- Time and space complexity are still typically exponential in the solution depth,  $d$  – i.e., the complexity is  $O(b^d)$  for some value  $b$ .
- However, A\* typically visits far fewer states than other optimal state-space search algorithms.

solution depth	iterative deepening	A* w/ Manhattan dist. heuristic
4	112	12
8	6384	25
12	364404	73
16	did not complete	211
20	did not complete	676

Source: Russell & Norvig, *Artificial Intelligence: A Modern Approach*, Chap. 4.

The numbers shown are the average number of search nodes visited in 100 randomly generated problems for each solution depth.

The searches do *not* appear to have excluded previously seen states.

- Memory usage can be a problem, but it's possible to address it.

## Implementing Informed Search

- Add new subclasses of the abstract Searcher class.

- For example:

```
public class GreedySearcher extends Searcher {
    private Heap<PQItem> nodePQueue;

    public void addNode(SearchNode node) {
        nodePQueue.insert(
            new PQItem(node, -1 * node.getCostToGoal ());
    }
    ...
}
```