


# State-Space Search

Computer Science E-22  
Harvard Extension School

David G. Sullivan, Ph.D.

## Solving Problems by Searching

- A wide range of problems can be formulated as *searches*.
  - more precisely, as the process of searching for a *sequence of actions* that take you from an *initial state* to a *goal state*



- Examples:
  - n-queens
    - initial state: an empty  $n \times n$  chessboard
    - actions (also called *operators*): place or remove a queen
    - goal state:  $n$  queens placed, with no two queens on the same row, column, or diagonal
  - map labeling, robot navigation, route finding, *many others*
- State space = all states reachable from the initial state by taking some sequence of actions.

## The Eight Puzzle

- A 3 x 3 grid with 8 sliding tiles and one “blank”

- Goal state:


	1	2
3	4	5
6	7	8

- Initial state: some other configuration of the tiles

- example:

3	1	2
4		5
6	7	8

- Slide tiles to reach the goal:




## Formulating a Search Problem

- Need to specify:
  1. the *initial state*
  2. the *operators*: actions that take you from one state to another
  3. a *goal test*: determines if a state is a goal state
    - if only one goal state, see if the current state matches it
    - the test may also be more complex:
      - n-queens: do we have n queens on the board without any two queens on the same row, column, or diagonal?
  4. the *costs* associated with applying a given operator
    - allow us to differentiate between solutions
    - example: allow us to prefer 8-puzzle solutions that involve fewer steps
    - can be 0 if all solutions are equally preferable

## Eight-Puzzle Formulation


- *initial state*: some configuration of the tiles
- *operators*: it's easier if we focus on the blank
  - get only four operators
    - move the blank up
    - move the blank down
    - move the blank left
    - move the blank right
- *goal test*: simple equality test, because there's only one goal
- *costs*:
  - cost of each action = 1
  - cost of a sequence of actions = the number of actions



	1	2
3	4	5
6	7	8

## Performing State-Space Search

- Basic idea:
  - If the initial state is a goal state, return it.
  - If not, apply the operators to generate all states that are one step from the initial state (its *successors*).




Consider the successors (and their successors...) until you find a goal state.

- Different search strategies consider the states in different orders.
  - they may use different data structures to store the states that have yet to be considered

## Search Nodes


- When we generate a state, we create an object called a *search node* that contains the following:
  - a representation of the state
  - a reference to the node containing the *predecessor*
  - the operator (i.e., the action) that led from the predecessor to this state
  - the number of steps from the initial state to this state
  - the cost of getting from the initial state to this state
  - an estimate of the cost remaining to reach the goal

```
public class SearchNode {
    private Object state;
    private SearchNode predecessor;
    private String operator;
    private int numSteps;
    private double costFromStart;
    private double costToGoal;
    ...
}
```



## State-Space Search Tree


- The predecessor references connect the search nodes, creating a data structure known as a *tree*.



- When we reach a goal, we trace up the tree to get the solution – i.e., the sequence of actions from the initial state to the goal.


## State-Space Search Tree (cont.)

- The top node is called the *root*. It holds the initial state.
- The predecessor references are the *edges* of the tree.
- depth* of a node  $N$  = # of edges on the path from  $N$  to the root
- All nodes at a depth  $i$  contain states that are  $i$  steps from the initial state:



## State-Space Search Tree (cont.)


- Different search strategies correspond to different ways of considering the nodes in the search tree.
- Examples:



## Representing a Search Strategy

- We'll use a *searcher* object.
- The searcher maintains a data structure containing the search nodes that we have yet to consider.
- Different search strategies have different searcher objects, which consider the search nodes in different orders.
- A searcher object may also have a *depth limit*, indicating that it will not consider search nodes beyond some depth.
- Every searcher must be able to do the following:
  - add a single node (or a list of nodes) to the collection of yet-to-be-considered nodes
  - indicate whether there are more nodes to be considered
  - return the next node to be considered
  - determine if a given node is at or beyond its depth limit

## A Hierarchy of Searcher Classes



- Searcher is an *abstract* superclass.
  - defines instance variables and methods used by all search algorithms
  - includes one or more *abstract methods* – i.e., the method header is specified, but not the method definition
    - these methods are defined in the subclasses
  - it *cannot* be instantiated
- Implement each search algorithm as a subclass of Searcher.

## An Abstract Class for Searchers

```
public abstract class Searcher {  
    private int depthLimit;  
  
    public abstract void addNode(SearchNode node);  
    public abstract void addNodes(List nodes);  
    public abstract boolean hasMoreNodes();  
    public abstract SearchNode nextNode();  
    ...  
    public void setDepthLimit(int limit) {  
        depthLimit = limit;  
    }  
    public boolean depthLimitReached(SearchNode node) {  
        ...  
    }  
    ...  
}
```

- Classes for specific search strategies will extend this class and implement the abstract methods.
- We use an abstract class instead of an interface, because an abstract class allows us to include instance variables and method definitions that are inherited by classes that extend it.

## Using Polymorphism

```
SearchNode findSolution(Searcher searcher, ...) {  
    numNodesVisited = 0;  
    maxDepthReached = 0;  
  
    searcher.addNode(makeFirstNode());  
    ...  
}
```

- The method used to find a solution takes a parameter of type Searcher.
- Because of polymorphism, we can pass in an object of *any* subclass of Searcher.
- Method calls made using the variable searcher will invoke the version of the method that is defined in the subclass to which the object belongs.
  - what is this called?


## Pseudocode for Finding a Solution

```

searcher.addNode(initialNode);


while (searcher.hasMoreNodes()) {
    N = searcher.nextNode();
    if (N is the goal)
        return N;
    if (!searcher.depthLimitReached(N))
        searcher.addNodes(list of N's successors);
}
    
```

- Note that we don't generate a node's successors if the node is at or beyond the searcher's depth limit.
- Also, when generating successors, we usually don't include states that we've already seen in the current path from the initial state (ex. at right).



## Breadth-First Search (BFS)

- When choosing a node from the collection of yet-to-be-considered nodes, always choose one of the shallowest ones.  
 consider all nodes at depth 0  
 consider all nodes at depth 1  
 ...
- The searcher for this strategy uses a *queue*.



```

public class BreadthFirstSearcher extends Searcher {
    private Queue<SearchNode> nodeQueue;

    public void addNode(SearchNode node) {
        nodeQueue.insert(node);
    }


    public SearchNode nextNode() {
        return nodeQueue.remove();
    }

    ...
}
    
```




## Tracing Breadth-First Search

search tree:




queue:



After considering all nodes at a depth of 1, BFS would move next to nodes with a depth of 2. *All previously considered nodes remain in the tree, because they have yet-to-be-considered successors.*

## Features of Breadth-First Search

- It is *complete*: if there is a solution, BFS will find it.
- For problems like the eight puzzle in which each operator has the same cost, BFS is *optimal*: it will find a minimal-cost solution.
  - it may *not* be optimal if different operators have different costs
- Time and space complexity:
  - assume each node has  $b$  successors in the worst case
  - finding a solution that is at a depth  $d$  in the search tree has a time *and* space complexity = ?



← 1 node at depth 0

←  $O(b)$  nodes at depth 1

←  $O(b^2)$  nodes at depth 2

- nodes considered (and stored) =  $1 + b + b^2 + \dots + b^d = ?$

## Features of Breadth-First Search (cont.)

- Exponential space complexity turns out to be a bigger problem than exponential time complexity.
- Time and memory usage when  $b = 10$ :


solution depth	nodes considered	time	memory
0	1	1 millisecond	100 bytes
4	11,111	11 seconds	1 megabyte
8	$10^8$	31 hours	11 gigabytes
10	$10^{10}$	128 days	<b>1 terabyte</b>
12	$10^{12}$	35 years	111 terabytes

- Try running our 8-puzzle solver on the initial state shown at right!

	8	7
6	5	4
3	2	1

## Depth-First Search (DFS)

- When choosing a node from the collection of yet-to-be-considered nodes, always choose one of the deepest ones.
  - keep going down a given path in the tree until you're stuck, and then backtrack
- What data structure should this searcher use?



```
public class DepthFirstSearcher extends Searcher {

    public void addNode(SearchNode node) {

    }

    public SearchNode nextNode() {


    }

    ...


}
```

## Tracing Depth-First Search (depth limit = 2)

search tree:




stack:




Once all of the successors of have been considered, there are no remaining references to it. Thus, the memory for this node will also be reclaimed.

## Tracing Depth-First Search (depth limit = 2)

search tree:




stack:



DFS would next consider paths that pass through its second successor. *At any point in time, only the nodes from a single path (along with their "siblings") are stored in memory.*

## Features of Depth-First Search

- Much better space complexity:
  - let  $m$  be the maximum depth of a node in the search tree
  - DFS only stores a single path in the tree at a given time – along with the “siblings” of each node on the path
  - space complexity =  $O(b \cdot m)$
- Time complexity: if there are many solutions, DFS can often find one quickly. However, worst-case time complexity =  $O(b^m)$ .
- Problem – it can get stuck going down the wrong path.
  - ➔ thus, it is neither complete nor optimal.
- Adding a depth limit helps to deal with long or even infinite paths, but how do you know what depth limit to use?




## Iterative Deepening Search (IDS)

- Eliminates the need to choose a depth limit
- Basic idea:
 

```

d = 0;
while (true) {
    perform DFS with a depth limit of d;
    d++;
}
      
```



- Combines the best of DFS and BFS:
  - at any point in time, we're performing DFS, so the space complexity is linear
  - we end up considering all nodes at each depth limit, so IDS is complete like BFS (and optimal when BFS is)

## Can't We Do Better?

- Yes!
- BFS, DFS, and IDS are all examples of *uninformed* search algorithms – they always consider the states in a certain order, without regard to how close a given state is to the goal.
- There exist other *informed* search algorithms that consider (estimates of) how close a state is from the goal when deciding which state to consider next.
- We'll come back to this topic once we've considered more data structures!
- For more on using state-space search to solve problems:  
*Artificial Intelligence: A Modern Approach.*  
Stuart Russell and Peter Norvig (Prentice-Hall).

- The code for the Eight-Puzzle solver is in code for this unit.
- To run the Eight-Puzzle solver:

```
j avac Ei ghtPuzzl e.j ava  
j ava Ei ghtPuzzl e
```

When it asks for the initial board, enter a string specifying the positions of the tiles, with 0 representing the blank.

example: for


	8	7
6	5	4
3	2	1

you would enter 087654321


- To get a valid initial state, take a known configuration of the tiles and swap *two pairs* of tiles. Example:

(you can also “move the blank” as you ordinarily would)

3	1	2
4		5
6	7	8



3	7	2
4		5
6	1	8



3	7	4
2		5
6	1	8